Machine Learning vs. Deep Learning

Muhammad Abdulfatah

Indiana Wesleyan University

3SU2026 Machine Learning Fundamentals (AIML-500-01N1U)

Dr. Longe

July 18, 2026

Machine Learning vs. Deep Learning: Real-World Applications

Machine learning and deep learning both allow computers to learn from data, but they solve problems in different ways. Traditional machine learning usually depends on people to select useful features before an algorithm is trained. Deep learning uses neural networks with many layers to discover useful features from raw data, especially images, audio, and language (Indiana Wesleyan University, n.d.). The distinction matters because a more complex model is not automatically a better model. The method has to fit the data, the decision being made, and the environment in which the system will operate.

For this report, I selected email spam detection as the traditional machine learning example and diabetic-retinopathy screening as the deep-learning example. Both are classification problems, yet the underlying data are very different. An email can be represented through known signals such as words, links, senders, and header information. A retinal photograph contains millions of pixels whose relationships are difficult to describe manually. Comparing these applications shows why traditional machine learning is suitable for one problem while deep learning is more appropriate for the other.

Machine Learning Example: Email Spam Detection

The course lesson uses email spam detection to illustrate traditional machine learning. A real-world implementation is Apache SpamAssassin, an open-source platform that system administrators use to classify and block unsolicited email. SpamAssassin does more than apply one model. It combines a scoring framework with tests on message headers and body text, Bayesian filtering, blocklists, and collaborative filtering databases (Apache Software Foundation, n.d.-a). Its sa-learn tool trains the Bayesian component from messages that users have labeled as spam or ham. The system then learns which words and other tokens are statistically associated with each category (Apache Software Foundation, n.d.-b).

Although the lesson names logistic regression and decision trees, SpamAssassin's Bayesian classifier follows the same traditional machine learning pattern. People supply labeled examples and define or expose meaningful features, and the model estimates the class of a new message. This approach is suitable because email provides many signals that can be represented directly: suspicious phrases, unusual punctuation, the number of links, sender reputation, domain information, message length, and header irregularities. These features are understandable to administrators and can be combined with explicit rules when the statistical model is uncertain.

Traditional machine learning also fits the operational setting. The model can be trained with a moderate collection of labeled messages, updated as users correct mistakes, and run efficiently on ordinary mail-server hardware. This matters because false positives can hide legitimate business messages. An administrator needs to investigate why a message was blocked and adjust a rule, allowlist, or training example. A conventional classifier offers more transparency and lower computational cost than a large neural network while still adapting to the organization's email flow.

Deep learning is not incapable of detecting spam. Large email providers can use deep language models to recognize context, obfuscated wording, and new attack patterns. However, it is less suitable for the specific SpamAssassin use case. A deep model would normally require more labeled data, specialized computing resources, a more complicated deployment process, and continuous monitoring for drift. Its decisions would also be harder for a local system administrator to explain. The possible gain in accuracy may not justify that additional cost and complexity. For an organization operating its own mail server, a hybrid of traditional statistical learning and understandable rules is the more proportionate solution.

Deep Learning Example: Diabetic-Retinopathy Image Recognition

For the deep-learning example, I selected image recognition in health care. Diabetic retinopathy damages blood vessels in the retina and can cause vision loss if serious disease is not identified and treated. Screening traditionally requires a trained professional to examine retinal photographs. Google and clinical partners developed an automated retinal disease assessment system, known as ARDA, to analyze those photographs and identify patients who should be referred for additional care. This is a direct application of the lesson's convolutional-neural-network image-recognition example.

The original development study trained a deep convolutional neural network with 128,175 retinal images that had been graded multiple times by a panel of eye-care professionals. When the model was evaluated on the separate EyePACS-1 and Messidor-2 datasets, the reported area under the receiver operating curve was 0.991 and 0.990, respectively (Gulshan et al., 2016). Those results showed that the model could distinguish referable disease with high sensitivity and specificity under controlled validation conditions.

The case later moved beyond a laboratory test. Aravind Eye Hospitals partnered with Google and Verily to use ARDA for diabetic-retinopathy screening in India. A 2025 post deployment study reported that the system had screened more than 600,000 patients across Tamil Nadu. In a randomly selected evaluation of 4,537 patients from 45 sites, ARDA achieved 97.0% sensitivity and 96.4% specificity for severe no proliferative or proliferative diabetic retinopathy. All patients with those severe forms were still referred for clinical care (Brant et al., 2025). This deployment demonstrates how deep learning can extend screening to communities where access to ophthalmologists is limited.

Deep learning is suitable because a retinal image is high-dimensional spatial data. Signs such as microaneurysms, hemorrhages, and exudates vary in size, position, color, contrast, and appearance. A convolutional neural network can learn this hierarchy directly from examples. Early layers recognize edges, shapes, and textures, while later layers combine those patterns into more complex representations associated with disease. The large expert-labeled image collection supports this data-intensive method, and the same model can evaluate photographs consistently once it has been validated.

Traditional machine learning is less suitable for the primary image task. Before training a conventional classifier, specialists and engineers would have to convert each photograph into handcrafted measurements such as lesion counts, vessel geometry, color statistics, or texture values. That process is labor-intensive and may discard subtle relationships among pixels. The features may also become fragile when camera models, lighting, image quality, or patient populations change. A traditional model could still combine the image result with structured information such as age or medical history, but it is less capable of learning the full visual complexity of the retina.

Deep learning's suitability does not remove the need for human oversight. Medical errors carry more serious consequences than an incorrectly filtered email. The system must be tested across patient groups and camera conditions, monitored after deployment, and designed to refer positive, uncertain, or ungradable cases to qualified professionals. The ARDA study is useful for this reason: it reports performance after deployment, not only during development. The model is valuable as part of a clinical workflow, but it should not be treated as a substitute for governance, quality control, or medical judgment.

Comparing the Two Approaches

The comparison is clearest when the data and constraints are placed side by side. Spam detection can use a manageable set of engineered features and must run quickly on modest infrastructure. Administrators also benefit from being able to explain and correct the system. Retinal-image recognition begins with raw visual data whose important patterns are difficult to define in advance. It has a large, labeled dataset, enough computing power to train a complex network, and a performance need that justifies the additional effort.

The choice therefore should not be framed as old technology versus new technology. It is a question of fit. Traditional machine learning is often the better choice for structured data, smaller datasets, limited computing resources, and situations where explanation is important. Deep learning is often the better choice for large collections of unstructured data and patterns that cannot be captured reliably through manual feature engineering. In both cases, leaders still have to define acceptable error rates, protect the training data, monitor performance, and decide when a person must review the model's output.

Conclusion

These cases made the distinction concrete. SpamAssassin shows why traditional machine learning remains useful: its features are understandable and it can run without extensive infrastructure. ARDA shows why deep learning is justified when complex images contain subtle patterns that must be learned from large expert-labeled datasets. My takeaway is that the best model fits the data, performs reliably, can be governed in practice, and supports accountable human decisions.

References

Apache Software Foundation. (n.d.-a). *Apache SpamAssassin*. https://spamassassin.apache.org/

Apache Software Foundation. (n.d.-b). *sa-learn: Train SpamAssassin's Bayesian classifier*. https://spamassassin.apache.org/full/4.0.x/doc/sa-learn.html

Brant, A., Singh, P., Yin, X., et al. (2025). Performance of a deep learning diabetic retinopathy algorithm in India. *JAMA Network Open, 8*(3), e250984. https://doi.org/10.1001/jamanetworkopen.2025.0984

Gulshan, V., Peng, L., Coram, M., et al. (2016). Development and validation of a deep learning algorithm for detection of diabetic retinopathy in retinal fundus photographs. *JAMA, 316*(22), 2402-2410. https://doi.org/10.1001/jama.2016.17216

Indiana Wesleyan University. (n.d.). *Machine learning vs. deep learning* [Course lesson].